

What do we talk about when we talk about Modules

Gašper Ažman (gasper.azman@gmail.com)
Andrew Bennieston (a.j.bennieston@gmail.com)

June 27, 2017

1 Aim of this paper

The Modules TS [1] has generated much discussion. This paper aims to capture the questions and decisions that are necessary to arrive at a coherent proposal.

This paper takes no stance on any of the issues highlighted. Where possible, the differing opinions are cited and summarised in the discussion of each section. It is likely that there are issues and discussions of which the authors are not aware. We do not claim to represent every facet of the design space, and expect further iteration for refinement.

The goal of the authors of this paper is to disentangle the various issues people have with modules, so that constructive argument can happen along the issues and requirements. The issues are presented in as orthogonal a manner as possible, so that loops in arguing are avoidable.

1.1 A note on terminology

This paper uses the word *name* in several contexts. It can be taken to mean *id-expression* as defined in [expr.prim.id] [2] or can be extended to include the names of macros. We attempt to call out the specific meaning in the cases where it is important. We also use the term *entity* as defined in [basic]. [2]

2 What Are Modules For?

This section presents all the various answers to the titular question that the authors are aware of. It is probable that not all the answers are mutually compatible, realistic, or even on topic. However, unless stated and subsequently addressed, they will keep surfacing.

Put another way, the subsections should be read in the context of *Should this be a design goal for modules?*

2.1 Tight control of visible names and entities

The various subsections deal with the visibility semantics of modules.

2.1.1 Control of exported names: Reduce the need for detail namespaces

Name lookup is a flexible and powerful feature of C++, but the rules are complex and it is no secret that the set of names (including names of macros) that are visible in user code is often larger than would be desirable. Customisation points further complicate matters.

To do this, private (and, increasingly often, public) features of libraries are implemented in a library-specific detail namespace, with externally visible names pulled into the public-facing main namespace with `using` directives.

This keeps implementations sane, but has the unfortunate side-effect of making compiler diagnostic messages longer than would be ideal.

One view is that it would be nice for modules to obviate the need to use detail namespaces to hide names by allowing name lookup outside the module to only find entities that are explicitly exported.

2.1.2 Control of imported names: Strict ordering guarantees for code inclusion

Another issue that keeps library authors awake at night is the introduction of new names into their namespaces due to code like

```
#include "otherlibrary.hpp"  
#include "mylibrary.hpp"
```

where `otherlibrary` defines some names (or macros) that clash with `mylibrary`.

An example of such a problem is the implementation of the standard library; all names in it need to use the `__name` and `_Name` conventions, despite already being in the `std` namespace – macros are not allowed to redefine names like that, since they are reserved. Regular libraries deal with this problem by assuming it won't happen.

If modules were specified to guarantee a strict partial ordering, such that *only the code that happened¹ above this module is the code it specifically imports* – possibly applied *non-transitively* for name visibility – then this issue disappears, and the standard library (along with every other library) can stop worrying about macros rewriting its code.

2.1.3 Transitivity of name visibility

Should `import moduleA`; also import the names from modules imported by `moduleA` transitively, or not?

Pro transitive The “better-precompiled-headers” approach dictates transitive name visibility. Namespaces remain the (only) method by which name visibility is managed, and modules are completely orthogonal.

Pro intransitive The existing practice from almost all other languages dictates non-transitivity of name visibility. This would mean explicit module imports everywhere, with allowance for re-exporting to enable *curated* library APIs.

2.2 Replacing precompiled headers

Modules offer an option for replacing precompiled headers. There are various reasons that precompiled headers are problematic, and adopting their solutions as design objective translates into certain choices. Therefore, we must know the problems with precompiled headers.

2.2.1 Compilation Speed: The *Parse It Once* Argument

C++ code takes a long time to parse, and longer to compile. There are many reasons for this, but having to read millions of lines of files to just get the declarations and definitions of all the visible, invisible, and irrelevant names, for *every compilation unit*, is one of them.

If there was a way to parse a header once, expose a data structure that would be able to load only the needed code at the time it is needed, and share that effort, the hope is that compilation times would drop.

2.2.2 The *Highlander* Argument

Precompiled headers are also very inflexible, since *there can only be one*, and therefore encourage having a pinch-point in the compilation graph that necessitates the recompilation of everything on every minor change of any header.

The hope is that modules would act as mix-and-match precompiled headers, keeping compilation quick while maintaining strict “as needed” import policies.

¹We use “happened” descriptively, because it means different things depending on how the other features of modules end up being specified. In the first instance, one could define it as simple textual inclusion of all (and *only*) transitive dependencies of a module. Since those could also be modules in their own right, the actual specification of this feature is difficult, and so we cautiously use the word “happened”.

2.3 The coupling between modules and namespaces

Some have the wish that modules would be better namespaces – modules would introduce namespaces, completing the C++-is-a-better-python dream².

This goal is complicated by the necessity for code (which may be modularised) to define entities in other namespaces. Standard library customisation points serve as a prime example; it is desirable for a module to be able to define a specialisation of e.g. `std::hash`, which suggests that modules and namespaces may need to be orthogonal.

Although the concepts may be orthogonal, many expect that there will be a strong correlation between modules and (outer) namespaces. The following ideas have been proposed to help with the common cases.

2.3.1 No coupling

This class of proposals sees modules and namespaces as completely orthogonal. Modules do not imply namespaces, they are file-based while namespaces are scoped, etc.

The problem with this approach is that while the two features are necessarily *independent* (to borrow from linear algebra), they are far from *orthogonal* when one considers how they are likely to be used. Not offering any support to help DRY³ can lead to confusion and poor ease-of-use.

2.3.2 Introduce a namespace by default and offer an opt-out.

A module can introduce a namespace by default, and then offer ways of “escaping”, for instance:

```
// All names implicitly in the foo::bar namespace and foo::bar module.
module foo::bar;

struct mystruct; // foo::bar::mystruct;

::std::hash<mystruct> {
    std::size_t operator(mystruct const&)
        const;
};
```

This style of opt-out was specified in a paper by Tristan Bridle. [3]⁴

2.3.3 Module-namespaces: Introduce a scope

A special syntax would enable introducing the module and its namespace together:

```
module namespace foo::bar {
    // All names in the foo::bar namespace and foo::bar module.
    struct mystruct;
} // end module foo::bar

module namespace foo::baz {
    // Look ma, two modules per file!
    struct mystruct;
}

// Entities defined here are only visible within the file and cannot be exported.
namespace std {
    struct std::hash<foo::bar::mystruct>
    {
        std::size_t operator(
            foo::bar::mystruct const&)
    }
```

²The authors make no claim as to whether this is a dream worth having.

³Don't Repeat Yourself

⁴The paper is not about modules.

```

    const;
};
} // namespace std

```

2.4 What do modules expose, and how?

Since modules seem to have the ambition of offering tight control of entities, they pose a question about what exactly they expose. The space of options follows. These are mostly mutually incompatible – macros are dealt with in a separate section.

2.4.1 Expose names

In this way, once you’ve exported a name, it’s visible, along with all its overloads (for functions) or specializations (for types).

2.4.2 Same as class specifiers

Have access specifiers `public`, `private` and `protected` work for modules the same way as they work for classes: expose names for the purposes of overload set resolution, but break if the resolved-to entity is not accessible (e.g. the name is private and the call is outside the module).

2.4.3 Expose entities

Only expose entities, explicitly. From outside the module, the private entities are hidden from all name lookup (i.e. non-exported function overloads do not appear in the overload set at all).

2.5 Are macros exported?

This is quite possibly the biggest current disagreement. This section summarises the arguments and counter-arguments.

2.5.1 Pro: Macros do appear in public interfaces of libraries

From `MAX_SIZE` to `assert` to `BOOST_HANA_STRING`, macros are a part of the public interface of many libraries, and therefore modules should offer a way to expose them, goes the argument.

A simple workaround is to supply the macro interface of the library in a separate header meant for the `include` mechanism. This header would also import the non-macro part of the interface it depends on. This makes the way modules would work with the preprocessor far more obvious.

However, this makes the general answer to *how do I use libraries* a lot *less* obvious; the answer may depend on whether you (or your clients) use the macros, or prefer a clean interface. Further, you may wish to consume the macros internally, but hide them from your clients.

The advantage is that libraries that are `imported` will *never* leak macros, which is a strong guarantee that is nice to have.

2.5.2 Contra (part 1): It breaks the compilation model

If macros are exported, then how can the import directive run after the preprocessor?

2.5.3 Contra (part 2): I want to guarantee my import doesn’t rewrite my source code

This is a valid concern, but can probably be solved independently of macro exports by looking at strict ordering guarantees for inclusion. It is assumed that you can trust the libraries you import explicitly.

2.6 A Better Pimpl

Modules should be able to be a compilation barrier. If only the (suitably defined) non-externally visible parts of a module change, a suitably mature implementation of a build system should not need to recompile the dependents.

This is mostly a quality-of-implemention issue.

3 Acknowledgements

We would like to thank Vittorio Romeo for a number of helpful suggestions, Jackie Kay for highlighting readability concerns and style suggestions, Tomas Puerle for the point about the better Pimpl, and Louis Dionne for proofreading the paper.

References

- [1] Gabriel Dos Reis. *Working Draft, Extensions to C++ for Modules (N4647)* <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2017/n4647.pdf>
- [2] ISO/IEC SC22 WG21 *Working Draft, Standard for Programming Language C++* <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2017/n4659.pdf>
- [3] Tristan Brindle. *Allowing Class Template Specialisations in Unrelated Namespaces* https://github.com/tcbrindle/specialization_proposal/blob/master/P0665r0.pdf